

Audible Feelings: Exploring the Role of Generative Sound in Reflective Mood Journaling

Ziheng Lyu
Whitman College
Walla Walla, Washington, USA

Jordan Wirfs-Brock
Whitman College
Walla Walla, Washington, USA




Figure 1: An overview of the *Audible Feelings* interaction flow: 1. The user receives a notification that their mood-logging window is open. 2. The user logs mood by inputting valence and intensity on sliders, and selecting emotion and influence tags. 3. The user responds, in text, to a question that encourages them to express their mood in metaphorical language. 4. A text-to-sound generative AI model generates an 8-second sound based on the users' text response. 5. The user must wait to hear the AI-generated sound, as it is available only after they log another mood entry the next day. 6. The user listens to the AI-generated sonic interpretation of the previous day's mood entry, and reacts to how well the audio matches. A full video walk-through is available in the supplementary materials.

Abstract

We present *Audible Feelings*, a mood-tracking application that uses generative AI to translate textual reflections into abstract soundscapes. While most mood-tracking tools rely on visual representations, in this work we explore the unique capabilities of sound for understanding emotional experience. Sound is immersive and often difficult to describe, inviting interpretive engagement rather than passive consumption. We applied a slow technology approach to foster anticipation and reflective distance through the temporal constraints of randomized logging windows and delayed listening. We situate this work as a design probe for the emerging space of sonic personal informatics. Drawing on reflections from our own use, as well as observations from a diary study, we share preliminary findings on the reflective possibilities of AI-generated sonic mood tracking and journaling. Our observations open onto a broader

question: What roles might sound play in personal journaling, and what does generative AI specifically enable within those roles?

CCS Concepts

• **Human-centered computing** → **Interaction design theory, concepts and paradigms.**

Keywords

Sound, Research through Design, Self-tracking, Slow Technology, Generative AI

ACM Reference Format:

Ziheng Lyu and Jordan Wirfs-Brock. 2026. Audible Feelings: Exploring the Role of Generative Sound in Reflective Mood Journaling. In *Adjunct Proceedings of the 14th Nordic Conference on Human-Computer Interaction (NordiCHI '26 Adjunct)*, October 03–07, 2026, Vaasa, Finland. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/3821402.3830094>

1 Introduction

Among personal informatics applications, one of the most common features is mood-tracking, whereby people can log how they are reacting to their emotions cognitively and physically [15]. Many digital mood-tracking applications prioritize quantification and



This work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.

NordiCHI '26 Adjunct, Vaasa, Finland

© 2026 Copyright held by the owner/author(s).

ACM ISBN 978-1-4503-XXXX-X/2018/06

<https://doi.org/10.1145/3821402.3830094>

visual pattern recognition, framing emotional experiences as discrete, comparable data points [3]. Through dashboards, they render mood as tidy, graphable, and notably *silent*. Such “quantified self” practices, while useful for identifying broad patterns, simplify the subjective, sensory, and temporal complexities of how people feel, interpret, and remember moods. Over-reliance on numerical data may cause feelings of anxiety, self-doubt, or alienation from our own bodies [8, 17]. In response, we ask: *How might we employ sound to create a more reflective and embodied engagement with mood-tracking practices?* We explore sound as a medium for evocative, open-ended interpretation because sound is temporal, immersive, and subjective. Unlike a visualization, a sound cannot be glanced at; it demands a different quality of attention. Sound must be experienced at its own pace: sound “imposes its duration on us” [1]. We embrace these sonic properties as design affordances that encourage reflection over recognition, and questioning over insight.

Even though most smartphones come with preinstalled voice memo apps, audio journaling is an uncommon practice. People are often more accustomed to documenting their lives by taking photos rather than gathering soundshots. Further, the ephemerality of sound means that, when people do notice an interesting sonic moment, by the time they open a voice memo app to record, the moment may have passed. One might instead select sounds from a pre-existing sound library that approximate a feeling, but such sounds carry no specific relationship to the personal experience that prompted them. Thus, instead of creating ways for people to capture or find sounds (which we pursue in other work), we focus on *deriving* sounds from the experiences and emotions people describe using text-to-sound generative AI. A generated sound is entangled with the user’s own language, produced from their specific input on a specific day. This shifts mood tracking beyond logging data to generating new *sonic* data, co-created with a generative AI system, that people can then feel and reinterpret.

To explore the reflective possibilities of AI-generated sound in mood-tracking practices, we developed *Audible Feelings* (see an overview in Fig. 1 and a video walk-through in our supplementary materials), a mobile mood-tracking and micro-journaling application that translates short, metaphorical text entries into 8-second audio clips. Rather than asking users to describe their emotions directly, the app prompts them to find a material or sensory metaphor [18] for their current state of mind (e.g., *What object best captures how you feel right now?*), which is then fed into the generative AI model along with a sonic label. We invite users to reflect on their emotional state through the gap between what they felt, how they described it, and what they hear. We also applied slow technology [6, 10, 11] as a design strategy to encourage reflection through anticipation through two temporal constraints: users can only log during a designated time-window each day, and cannot listen to the sonic interpretation of their mood entries until after they have logged the next day.

In this ongoing research-through-design [19] work, we treat *Audible Feelings* as a reflective technology probe [7, 16] that allows us to understand the tensions people experience when generative audio is incorporated into mood-tracking, refine the system in situ, and generate new design directions and questions. As researchers and designers, we also draw on autobiographical design

traditions [9], using the app ourselves during the design process. In this work-in-progress paper, we describe the key design decisions behind *Audible Feelings*, share observations that have emerged from our first-person engagement with the system over the past 10-months of development, and describe preliminary observations from a two-week diary study with $n=7$ participants. We contribute preliminary findings on the reflective possibilities of AI-generated sonic mood tracking, and use this work-in-progress to invite feedback on the design space in understanding how AI-generated sound might support reflective mood tracking.

2 Background

2.1 Reflective personal informatics

With the rise of digital journaling and personal informatics, self-tracking practices have become increasingly data-driven, emphasizing systematic logging, visualization, and pattern recognition across domains such as mood, health, and productivity [14]. While self-understanding and reflection have been implicit goals in self-tracking apps, Cho et al. note that apps often fall short of truly supporting reflection [2]. To better support reflection, they suggest that designers emphasize qualitative journaling, empower users to customize the data collection process, and present a blend of system- and user-created data representations [2]. Recent advancements in generative AI open up new possibilities for creating reflective personal informatics technologies, especially in the area of co-authored data representations. For example, Soobin Park et al. designed a system that generates AI images based on self-tracked data, demonstrating how AI-generated media can be a tool for reflecting on personal data [13]. We extend this tradition by investigating how people might reinterpret their mood data by reflecting on AI-generated sound.

2.2 Sonic ways of knowing in personal journaling

Sound is a promising medium for reflection because we feel sound viscerally in our bodies. Philosopher Christoph Cox notes that sounds are vibrations that move through matter, “as material as sticks and stones...[they] set eardrums aquiver, rattle walls, and shatter wine glasses” [5]. Tying this to principles of reflective design [16], sound crosses a boundary—between digital device and human body—as well as subverting dominant visual metaphors for representing data. Recent HCI research has explored the reflective possibilities of AI-generated sounds in the context of journaling practices. Joonyoung Park et al. have developed *NoRe* [12], an application that uses generative AI to create music based on long-form text journal entries. And Choi et al. developed *SoulNote* [4], a tool for deaf and hard-of-hearing users to generate music based on a conversation-like journaling interaction. Both works demonstrate that AI-generated music is a promising technology for supporting emotional exploration. We extend this theme by exploring *non-musical* AI-generated sounds, such as ambient soundscapes and sound effects, in addition to music. We also use abstract metaphorical questions that can elicit sensory-rich language from users that is specifically meant to produce interesting sonic translations [18].

3 Designing Audible Feelings

We developed *Audible Feelings* through an iterative RtD process [19], building and using the app ourselves over an extended period and reflecting on what we observed along the way. First, we spent several weeks in open-ended exploration of the ElevenLabs text-to-sound model¹, feeding it a range of inputs to explore the capabilities of AI-generated sound and its potential applications in mood tracking and micro-journaling. We tried inputs that seemed straightforward to represent sonically (e.g. “A memorable car ride”) and inputs that resembled language people might use when reflecting on their day (e.g. “Today I feel relaxed, maybe a little too relaxed, but also very motivated at the same time.”) We noticed that the AI often got overwhelmed when the text input was long and detailed. So we focused on short, metaphorical text inputs that could produce a range of possible sounds. We then created *Audible Feelings* as a mobile app that allowed users to log their mood as well as enter a short text note that would be used as a prompt for the ElevenLabs API to create a sound. We intended *Audible Feelings* to be a probe for exploring the reflective possibilities of AI-generated sound, thus we made decisions that supported openness and flexibility. We present our key design decisions below, not as finalized solutions, but as a series of explorations that can open up further design conversations.

3.1 Crafting metaphorical questions

We noticed that direct, emotional language tended to produce outputs that felt predictable, collapsing the gap between what was written and what was heard in ways that left little room for reflection. Words like “anxious,” “happy,” or “crying” guided the model toward outputs that illustrated the emotion in recognizable, culturally familiar terms—bright major-key tones for “happy,” a literal crying sound for “crying”—producing sounds that mirrored what the user already felt rather than opening it up for reinterpretation. We were drawn to a different quality of encounter, one where the relationship between the text input and the generated sound was less immediately legible, requiring the user to actively make sense of it. Instead, we began experimenting with questions that invited users to find a sensory analogy for their current state. Questions like “*What plant or flower represents how you’re feeling?*” or “*If your mood were a time of day, what would it be?*” invited users to translate their internal states into material references before these were passed to the generative audio model (a full list of questions is available in our supplemental materials). A response like “*The growing beautiful tulips in my yard*” or “*A late afternoon with nothing left to deal with*” gave the generative AI model something less straightforward. The outputs became more varied and more surprising as a result, introducing some protective distance from raw emotional language.

3.2 Introducing sonic labels

As we developed and lived with the app, we encountered outputs from the generative sound model that we had not anticipated, including some that sounded close to vocalized speech. While we found these interesting from a research standpoint (as we discuss in section 4.2), to support our goals of sense-making and reflection, we wanted to avoid cases where the generated sound attempted to

directly “read” the text input. Thus, we began looking for ways to shape the character of the generated sound. Adding a short sonic label prepended to the text input before it was sent to the API made a meaningful difference. We experimented with several options including “ambient soundscape,” “piano solo,” “orchestral,” “jazz music,” “acoustic guitar,” and, crucially, no label at all. We found that each label oriented the model toward a distinctive sonic territory. In the current version of the app, a label is assigned randomly for each entry. We want to be clear that this is an exploratory decision that supports our goals to use *Audible Feelings* as a probe: random labels allow us to keep the outputs unpredictable without being shaped by user expectations or our own assumptions about what would sound right. Through using the app ourselves, we began to notice patterns in how different labels shaped the outputs. Some musical genre labels tended to produce outputs that felt strongly genre-typical: A response prepended with “orchestral” might yield a generic orchestral piece that has little apparent relationship to what was written. In contrast, “no label” produced outputs that were less generic and sometimes more surprising. A generation prepended with a non-musical label, like “ambient soundscape,” sits somewhere in between.

3.3 Building in temporal friction

Drawing on slow technology principles to encourage anticipation and reflective distance [6, 10, 11], we introduced two forms of deliberate temporal friction: a time-bounded logging window, and a mandatory one-day delay before listening to generated sounds. Mood logging is available only during a two-hour window, either at a fixed time the user sets or a randomized point within a longer range. The fixed window can build routine and consistency, while the randomized window introduces a degree of serendipity, catching mood at unexpected moments throughout the day. Once a mood entry is logged, the generated sound is locked until the following day’s entry is completed, creating a mandatory one-day delay between logging and listening. We designed this delay to invite users to return to what they made the day before with some reflective distance from the emotional state that generated it. Interacting with generative AI typically involves a moment of waiting, a few seconds to a few minutes, before the model returns an output. We have come to expect this pause as part of the interaction rhythm. *Audible Feelings* plays on this paradigm, extending the delay to a full day, in order to reimagine temporality in human-AI interaction. In practice, we found ourselves genuinely curious about what the previous day’s sound would be like, which motivated us to be more consistent in our use of the app.

4 Living with Audible Feelings

The authors have been developing and using *Audible Feelings* intermittently over the course of 10 months. Below, we reflect on our own relationships with the app, tracing how particular sonic encounters shaped our understanding of what the system does and what it might become. We present these reflections as first-person vignettes to preserve their autobiographical character.

¹<http://elevenlabs.io/app/sound-effects>

4.1 Ziheng Lyu's Reflection: Beyond Concordance

One sound has stayed with me longer than most. The question was, "If your mood were a time of day, what would it be?" and I responded with a single word: evening, because I was feeling a bit tired and just craving some alone time, and evening is usually associated with alone time to me. No sonic label was attached, so the API received nothing more than that. What came back was a quiet ambient soundscape with very distant, light raindrops and the occasional sound of a car passing by. It felt almost disarmingly right.

I often play white noise or the hum of a small fan to fall asleep to, and I enjoy waking up slowly on unhurried mornings to rain pattering. This generated sound fit that atmosphere so precisely that it felt less like a translation of my mood and more like the system had produced something I already knew I loved. I understood roughly why "evening" might yield something quiet. But the distant rain and the passing car are sounds that I have trouble accounting for with just a single word, "evening." I did not journal "a rainy night" or "distant city street ambience at dusk" to specify I wanted to perhaps hear some rain sounds or the sound of a car passing by.

I am surprised by my own reaction to the sound. Usually when a generated sound feels concordant, my response is a simple recognition: "yes, it got it," and I move on. This one stopped me because it was a bit too concordant, too close to something genuinely personal, and that excess of fit made me want to understand it rather than just accept it. I relistened to this sound many times to interrogate it: in what specific ways did this sound feel right, and what did that reveal about what "evening" meant to me beyond its literal time of day?

That experience shifted how I approach listening in the app more broadly. I began paying more attention to sounds that feel concordant, not just ones that surprise or confuse me. It made me realize that a sound can be "confusingly" concordant — that excess of fit raises as many questions as a mismatch would. In the weeks that followed, I found myself returning to concordant sounds with new questions: not "did the AI get it right," but "what does it mean that I experience it as right, and what does that reveal about my own sonic associations and emotional habits." This realization has since shaped how I understand what the app is doing, which is less about producing accurate representations of mood and more about producing occasions for self-inquiry.

4.2 Jordan Wirfs-Brock's Reflection: Probing the Boundaries of Speech

I have always been interested in where the line between speech and non-speech lies, where a groan becomes an expression, in what meaning is imbued in our diction and tone that words alone cannot capture. Thus, my favorite generated sound is one of a human-like voice singing. In response to the question, "Share a food that encapsulates your feeling," I wrote, "Oooh, I am feeling the matcha chocolate chip cookies I made today. Confusing but satisfying." And the sound, I will try to describe it although it feels indescribable when I listen: It starts with a male-sounding voice singing a swinging, upbeat melody that reminds me of the mid-twentieth century, although I can't quite pinpoint a genre (the tag the system appended was "Jazz music"). The song has lyrics, but only some of them are words. The voice

starts, "Ooooh-oooh-oh-oooh," in a descending glissando, then transitions smoothly to, "T-ah-hai-cam-poko-cookies I am made today." It is mostly acapella, with the exception of one short horn-ish blast underscoring the first syllable of "cookies," followed by something that sounds like shaking a jar of rice.

It is such a strange sound, with a humor and a lightness to it. I find joy in it. I have listened to it on loop, probably hundreds of times. I have hummed along to it, and I have tried to sing it (with limited success). I appreciate how it clearly draws inspiration from the words I wrote, but is not bound to those words. The sound is not trying to recreate my memory of that day — which the system has no real access to — but is an augmentation to that day that, through its whimsy and strangeness, invites me to return to my memories on my own terms.

My reaction to this sound contrasts starkly with how I encountered generated sounds that resembled human voices on previous occasions, when they felt invasive, and like they were exposing the cracks of generative audio. For example, in response to my text entry, "Feels like I am playing life on hard mode — but that is kind of just how parenting is. Dealing with some sick toddlers, oof," I got a generated sound of a female voice speaking a purely gibberish language. And yet I try to transcribe it, something like, "Belo ray how nay bot, your sad culture newz ainshake an hans mar." I pick out some English words, but am not sure if they are really there, or if I am straining to hear them through the nonsense. On a different day, a different female voice read back my words precisely, "Had a stressful morning, realized that I have too much to do." But that wasn't enough: the system then appended a mysterious whisper (now a different voice) saying, "lavos twenty-three." Directly reading my words was disorienting and intrusive — and not at all what I wanted — but the addendum felt like an insult, like my words weren't enough.

There is something creepy about the system attempting, but failing, to turn my words into human speech. There is no song there, no music. Yet there is something delightful when the system riffs on my words, turning them into a song. I still don't know exactly where the line is, when computer-generated human speech transitions from sinister to playful, but the app is helping me probe that space.

5 Observations from a Diary Study

To situate our autobiographical reflections in a broader context, we draw on preliminary observations from a diary study of sonic mood-tracking and journaling (n=7). In this work in progress, we offer a brief overview of our methods: Participants completed an initial 30-minute interview about their mood-tracking and journaling practices, then used the app for two weeks, leaving voice-memo reflections every few days. At the end of the study, they participated in a 90-minute co-design workshop. Six participants identify as female and one as non-binary. The participants are between ages 18 and 54, are students and staff at an institution of higher education, and represent a range of geographic backgrounds, including Colorado, California, Peru, Indonesia, Malawi, Tajikistan, and Switzerland. All participants had a prior relationship with journaling, though for three it was a past rather than current practice. Three participants had significant musical backgrounds, including singing and music technology. This study was approved by our institution's internal review board.

As a complement to our autobiographical reflections, in this section we present two themes that emerged from early observations of the participant study: generative sound as impressionistic memory, and playfulness to support sensory reflection.

5.1 Generative Sound as Impressionistic Memory

As they used *Audible Feelings*, participants treated sound as a prompt for recall rather than as a form of direct documentation. Two participants noted that the generated sound had the potential to capture the overall mood of a day; one articulated this by noting that the generated sounds would often match “the vibe of the day but not the detail.” The sound captured something about the felt quality of their experience without representing its specifics. Rather than experiencing this as a limitation, participants found this generative; the sound returned them to an atmosphere of their day that they could then narrate themselves. This impressionistic quality resonated with another participant, who had previously tried several journaling apps but found them too demanding of detailed input to be sustainable. They described the metaphorical questions and short-text responses of *Audible Feelings* as a low-barrier way of reflecting on the day that removed the pressure of thorough documentation.

This impressionistic quality is built into the design of *Audible Feelings* at two levels: First, the metaphorical questions are intentionally designed to steer users away from direct emotional description. Rather than asking how they feel, the questions ask what object, environment, texture, or creature their feeling resembles. This framing invites a form of oblique self-expression, where the act of finding a metaphor is itself an interpretive gesture. Second, the AI sound generation takes this already-impressionistic input and transforms it further, producing an audio artifact shaped by the system’s own learned associations between language and sound. Users encounter something that resonates with, departs from, or refracts their original input in ways that invite active sense-making.

5.2 Playfulness to Support Sensory Reflection

Several participants noted that the metaphorical questions and impressionistic sounds felt playful in a way that distinguished them from other journaling formats. One participant described the app as feeling “lighthearted,” adding that the questions had a quality of “silliness” that made the act of logging feel less weighty and more inviting. Another participant noted that, on days when they were not in the headspace to reflect deeply, this playful quality made the app approachable because it still gave them a way to mark how they felt without demanding a taxing process of self-examination. Further, receiving an AI-generated sound offers users a whimsical artifact to reflect upon which may feel more comfortable than reflecting on their mood entries directly.

As the participants described their experience with *Audible Feelings*, they directly ascribed this playfulness to the metaphorical question format. Questions like “Describe your mood as a texture,” or “If you were to be an animal today, what would it be?” invite a kind of imaginative, low-stakes engagement that differs from conventional journaling prompts. Finding an unexpected analogy for one’s emotional state has a lightness to it, and this quality makes space

for reflection even when it does not feel immediately accessible. Although we implemented this design feature with the intention of inspiring interesting abstract sounds, we note that this kind of metaphorical prompting is a promising technique for other kinds of cross-sensory journaling.

6 Discussion and Future Work

We began this work by asking, *How might sound create a more reflective and embodied engagement with mood tracking?* Our design process and preliminary observations suggest that this question opens onto a broader one: *What roles might sound play in personal journaling, and what does generative AI specifically enable within those roles?* Through building and living with *Audible Feelings*, we have begun to identify several possible roles for sound in journaling. Sound can serve as a **tool for documentation**, capturing a moment or atmosphere as it happens. It can also function as a **memory trigger**, returning a listener to a felt experience from the past. It can be a medium for **processing emotions**, giving affective states a form that can be listened to and sat with. It can serve as **creative material**, inviting a playful and imaginative engagement with one’s inner life. And, as our findings begin to suggest, it can function as an **impressionistic distillation** of a day, capturing its overall felt quality without attempting a thorough account. Not all of these roles are equally developed in this paper, and we see mapping this design space as an important direction for future work.

While some of the roles outlined above can be met through recording or finding sounds, we suggest that generative AI opens up three distinct possibilities that neither recording nor finding can offer. First, it can **fill sonic gaps**, producing sounds for experiences that were never captured—or would be impossible to capture—in the moment. Second, it can **generate sounds for things that have no real-world sonic referent**, like, “a frustrating day that resembles a cactus.” Third, and most directly suggested by our preliminary findings, it can **interpret and distill**: taking a user’s own metaphorical language and transforming it into something evocative rather than accurate, producing an artifact that is entangled with the user’s experience while remaining open to reinterpretation. Our explorations into metaphorical prompts and sonic labels reveal how the interpretive gap between input and output can be calibrated to support reflection. We plan to continue this work to better understand how to use generative sound without tipping into frustration or irrelevance.

Our design exploration suggested several other directions for future work: Perhaps AI-generated mood sounds might function as affective translators, making inner states more legible to others in ways that text or numbers cannot. We also hope to explore how extended timescales and long-term reflection might reshape the relationship between sound and the experiences they index. A sound that feels only loosely connected to a day at the time of listening might resonate differently months later, when the emotional context has shifted. What might happen if users are able to listen to a week’s worth of generated sounds, or a month’s worth, hearing a sonic distillation of a season? Another direction is to explore how we might incorporate user-recorded sounds alongside AI-generated ones, creating a reflective dialogue between experienced and AI-generated sonic representations.

7 Conclusion

We presented *Audible Feelings*, which serves as a technology probe to examine the role of generative AI sounds in mood-tracking and journaling practices. We demonstrated how metaphorical questions can serve dual design purposes of supporting playful yet meaningful self-reflection, as well as providing the sensory openness needed to produce AI-generated sounds that can provoke sense-making. We suggest future work that examine the distinct roles sound can play in mood-tracking and journaling practices, and the ways that generative audio can fill those roles.

Acknowledgements

This work is supported by NSF grant #2348244. We thank Catelyn Clevenger for assisting with the diary study, as well as the members of the Resonate Data Lab who gave valuable feedback throughout this work.

References

- [1] Michel Chion. 2015. *Sound: An acoulogical treatise*. Duke University Press.
- [2] Janghee Cho, Tian Xu, Abigail Zimmermann-Niefield, and Stephen Volda. 2022. Reflection in Theory and Reflection in Practice: An Exploration of the Gaps in Reflection Support among Personal Informatics Apps. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems* (New Orleans, LA, USA) (CHI '22). Association for Computing Machinery, New York, NY, USA, Article 142, 23 pages. doi:10.1145/3491102.3501991
- [3] Eun Kyoung Choe, Nicole B Lee, Bongshin Lee, Wanda Pratt, and Julie A Kientz. 2014. Understanding quantified-selfers' practices in collecting and exploring personal data. In *Proceedings of the SIGCHI conference on human factors in computing systems*. 1143–1152.
- [4] Youjin Choi, JinYoung Yoo, JaeYoung Moon, Yoonjae Kim, Eun Young Lee, Jennifer G Kim, and Jin-Hyuk Hong. 2026. From Daily Song to Daily Self: Supporting Emotional Growth of Deaf and Hard-of-Hearing Individuals through Generative AI Songwriting. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems*. 1–26.
- [5] Christoph Cox. 2017. Sonic thought. *Sonic thinking: A media philosophical approach* (2017), 99–109.
- [6] Lars Hallnäs and Johan Redström. 2001. Slow technology—designing for reflection. *Personal and ubiquitous computing* 5, 3 (2001), 201–212.
- [7] Hilary Hutchinson, Wendy Mackay, Bo Westerlund, Benjamin B. Bederson, Alison Druin, Catherine Plaisant, Michel Beaudouin-Lafon, Stéphane Conversy, Helen Evans, Heiko Hansen, Nicolas Roussel, and Björn Eiderbäck. 2003. Technology probes: inspiring design for and with families. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems* (Ft. Lauderdale, Florida, USA) (CHI '03). Association for Computing Machinery, New York, NY, USA, 17–24. doi:10.1145/642611.642616
- [8] Deborah Lupton. 2013. Quantifying the body: monitoring and measuring health in the age of mHealth technologies. *Critical public health* 23, 4 (2013), 393–403.
- [9] Carman Neustaedter and Phoebe Sengers. 2012. Autobiographical design in HCI research: designing and learning through use-it-yourself. In *Proceedings of the designing interactive systems conference*. 514–523.
- [10] William Odom. 2024. Illustrating, annotating & extending design qualities of slow technology. In *Adjunct Proceedings of the 2024 Nordic Conference on Human-Computer Interaction*. 1–1.
- [11] William Odom, Erik Stolterman, and Amy Yo Sue Chen. 2022. Extending a theory of slow technology for design through artifact analysis. *Human-Computer Interaction* 37, 2 (2022), 150–179.
- [12] Joonyoung Park, Hyewon Cho, Hyehyun Chu, YeEun Lee, and Hajin Lim. 2025. NoRe: Augmenting Journaling Experience with Generative AI for Music Creation. In *Proceedings of the 2025 ACM Designing Interactive Systems Conference (DIS '25)*. Association for Computing Machinery, New York, NY, USA, 2718–2737. doi:10.1145/3715336.3735845
- [13] Soobin Park, Hankyung Kim, and Youn-kyung Lim. 2025. Reimagining Personal Data: Unlocking the Potential of AI-Generated Images in Personal Data Meaning-Making. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. 1–25.
- [14] John Rooksby, Mattias Rost, Alistair Morrison, and Matthew Chalmers. 2014. Personal tracking as lived informatics. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems* (Toronto, Ontario, Canada) (CHI '14). Association for Computing Machinery, New York, NY, USA, 1163–1172. doi:10.1145/2556288.2557039
- [15] Stephen M Schueller, Martha Neary, Jocelyn Lai, and Daniel A Epstein. 2021. Understanding People's Use of and Perspectives on Mood-Tracking Apps: Interview Study. *JMIR Ment Health* 8, 8 (Aug. 2021), e29368. doi:10.2196/29368
- [16] Phoebe Sengers, Kirsten Boehner, Shay David, and Joseph 'Jofish' Kaye. 2005. Reflective design. In *Proceedings of the 4th decennial conference on Critical computing: between sense and sensibility*. 49–58.
- [17] Tamar Sharon. 2017. Self-tracking for health and the quantified self: Re-articulating autonomy, solidarity, and authenticity in an age of personalized healthcare. *Philosophy & Technology* 30, 1 (2017), 93–121.
- [18] Jordan Wirfs-Brock, Maxene Graze, Laura Devendorf, Audrey Desjardins, Visda Goudarzi, Mikhaila Friske, and Brian C Keegan. 2022. Sketching across the senses: Exploring sensory translation as a generative practice for designing data representations. In *CHI Conference on Human Factors in Computing Systems Extended Abstracts*. 1–7.
- [19] John Zimmerman, Jodi Forlizzi, and Shelley Evenson. 2007. Research through design as a method for interaction design research in HCI. In *Proceedings of the SIGCHI conference on Human factors in computing systems*. 493–502.